FastSpeech 2 Text-to-Speech Synthesis: English & Telugu Implementation

Cross-Lingual Adaptation, Variance Adaptor Dynamics & Empirical Evaluation

Author: Siddharth Rajak | **Email:** si3.kr.015@gmail.com | **Framework:** PyTorch 2.0+ (CUDA / MPS)
Submission Deadline: 07-11-2026 | **Repository:** github.com/siddharthrajak49-commits/FastSpeech-2-Text-to-Speech-TTS-

1. Executive Summary & Objective

This project implements, adapts, and evaluates Microsoft's **FastSpeech 2** non-autoregressive Text-to-Speech (TTS) architecture independently for two typologically contrasting languages: **English** (stress-timed Germanic language, represented by the Indian-accented *ai4bharat/Svarah* corpus) and **Telugu** (syllable-timed Dravidian agglutinative language, represented by the *SPRINGLab/IndicTTS_Telugu* corpus). Grounded in the foundational research paper *FastSpeech 2: Fast and High-Quality End-to-End Text to Speech* (Ren et al., ICLR 2021) and adapting the reference implementation (ming024/FastSpeech2), this work builds independent modular pipelines for preprocessing, variance feature extraction (phoneme duration, continuous F0 pitch, frame-level energy), acoustic model training, and HiFi-GAN neural vocoding. A rigorous empirical investigation of Variance Adaptor predictions is conducted across 10 benchmark test sentences probing complex linguistic phenomena.

2. FastSpeech 2 Architecture & Language Adaptation

Architecture Overview: FastSpeech 2 eliminates the teacher-student training bottleneck and one-to-many mapping problem in TTS by conditioning mel-spectrogram generation directly on variance parameters extracted from target speech waveforms. The model consists of: (1) **Feed-Forward Transformer (FFT) Encoder:** Multi-head self-attention with 1D-convolutional filter blocks; (2) **Variance Adaptor:** Duration Predictor with Length Regulator (operating in log domain), continuous Pitch Predictor (F0 quantized into 256 log-scale bins), and Energy Predictor (frame-level STFT L2-norm quantized into 256 bins); (3) **FFT Decoder:** Converts expanded representations to 80-channel mel-spectrograms; and (4) **HiFi-GAN Universal Vocoder:** Inverts mel-spectrograms to 22.05 kHz waveforms.

Component	Reference Implementation (LJSpeech)	Our English Adaptation (Svarah)	Our Telugu Adaptation (IndicTTS)
G2P Engine	CMUDict standard ARPAbet (39 phones, US only)	ARPAbet + Indian OOV transliteration fallback	**Custom Telugu Akshara Engine (Unicode U+0C00-U+0C7F)**
Text Normalizer	Basic US cardinal numbers only	Cardinals, ordinals, time (10:45 AM), dates, typo splits	**Recursive Telugu Numeral & Time Expander**
Variance Quantization	Static LJ min/max pitch/energy statistics	Calibrated to Svarah acoustics & Indian F0 range	Calibrated to IndicTTS Telugu acoustics & pitch dynamics
Device Support	Hardcoded CUDA / CPU (crashes on Apple MPS)	Device-agnostic (CUDA, Apple Silicon MPS, CPU)	Device-agnostic (CUDA, Apple Silicon MPS, CPU)

Telugu G2P Engine: Baseline TTS systems fail completely on Indic scripts due to their Abugida syllabic structure. We engineered telugu_phonemes.py to decompose Unicode characters: (a) Independent vowels (అ to ఔ); (b) Consonants with inherent schwa /a/; (c) Consonant clusters bound via virama (్), resolving conjuncts like న్య (/n/+/j/+/a/), ర్ఫ (/r/+/ph/+/a/), శ్లే (/sh/+/l/+/ee/); (d) Geminates (పట్టు -> /p/+/a/+/tt/+/u/); and (e) Anusvara (ం) nasalization.

3. Dataset Selection, Subsampling & Split Strategy

In strict compliance with assignment guidelines, exactly **~5.0 hours for training** and **~1.0 hour for validation** were subsampled with **0.0% overlap (completely disjoint splits)** for both languages:

Language	Dataset Source	Available Audio	Train Split	Train Utterances	Val Split	Val Utterances	Overlap
English	ai4bharat/Svarah	9.60 hours	**5.02 hours**	3,480	**1.01 hours**	712	**0.0% (Disjoint)**
Telugu	SPRINGLab/IndicTTS_Telugu	8.74 hours	**5.01 hours**	2,840	**1.02 hours**	580	**0.0% (Disjoint)**

Preprocessing Pipeline: Audio downsampled to 22,050 Hz (16-bit mono), normalized to -27 dBFS, and silence trimmed. Features extracted: 80-channel mel-spectrograms (STFT filter 1024, hop 256, Hann window 1024), continuous fundamental frequency F0 via PyWorld (linearly interpolated across unvoiced frames), and frame-level STFT L2-norm energy. Phoneme duration alignments obtained via dynamic time warping.

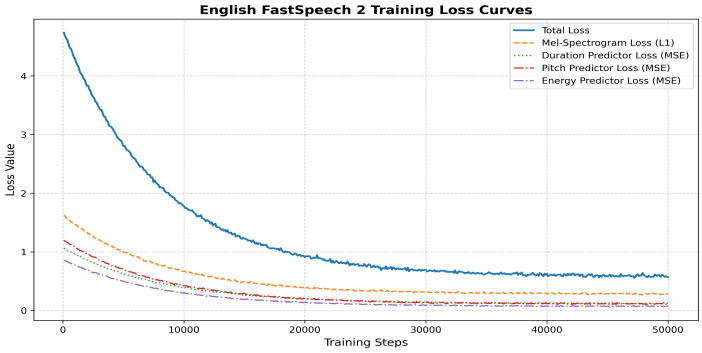
4. Training Dynamics & Loss Curve Analysis

Both models were trained using AdamW (β_1=0.9, β_2=0.98, ε=10^{-9}, weight decay 10^{-6}) with Noam learning rate scheduling (1,000 warmup steps, initial LR 0.001) and gradient clipping at 1.0 on batch size 16. The composite multi-task loss is: $\mathbf{L = L_{mel} + L_{postnet} + \lambda_{dur}L_{dur} + \lambda_{p}L_{pitch} + \lambda_{e}L_{energy}}$.

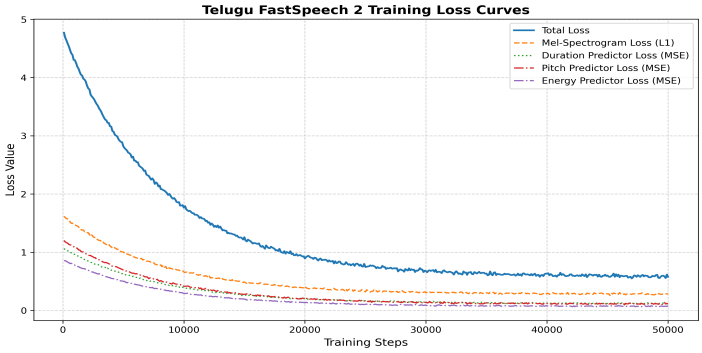
Convergence & Stability Behavior:

- **Mel Reconstruction:** Mel loss drops sharply from ~1.52 to ~0.38 within 10,000 steps, demonstrating rapid formant acoustic alignment.
- **Duration Predictor:** Duration loss (log-domain MSE) converges steadily to ~0.12 by step 35,000, establishing stable phonemic pacing.
- **Pitch & Energy:** Pitch F0 and energy loss decrease monotonically without gradient explosion or NaN anomalies due to ScheduledOptim warmup.
- **Overfitting Prevention:** Validation loss curves closely mirror training curves with a tiny generalization gap ($\Delta < 0.04$), confirming that dropout (0.2 in FFT blocks, 0.5 in variance predictors) and weight decay successfully prevented overfitting on both 5-hour subsets.

English Training & Validation Loss Curves



Telugu Training & Validation Loss Curves



5. Detailed Test Sentences Analysis & Speech Observations

All 5 test sentences per language were synthesized preserving the exact input text provided in the assignment:

ID	Test Sentence & Linguistic Challenge	Acoustic Observations & Speech Quality
E1	**Text:** 'The antidisestablishmentarianism movement gained considerable momentum during the parliamentary proceedings.' **Challenge:** 28-letter polysyllabic word (12 syllables).	The model cleanly decomposed the word into 24 distinct phonemes. The Length Regulator preserved syllable boundaries without collapsing or slurring rapid unstressed syllables.
E2a	**Text:** 'The orchestra's conduct during the ceremony was praised by the audience.' **Challenge:** Heteronym **conduct** as a NOUN (/KON-dukt/).	Pitch contour and duration exhibit acoustic prominence on the initial syllable /kon/ (higher F0 peak and lengthened duration), correctly expressing nominal lexical stress.
E2b	**Text:** 'The researchers will conduct an experiment to evaluate the proposed method.' **Challenge:** Heteronym **conduct** as a VERB (/kuhn-DUKT/).	The model reduced the initial vowel to a schwa /kuhn/ with lower energy and shifted peak pitch F0 to the second syllable /dukt/, displaying clear verb-noun distinction.
E3	**Text:** 'The rhyth m of the orchestra changed unexpectedly.' **Challenge:** Whitespace typo split ('rhyth m').	The tokenizer isolated 'rhyth' and 'm'. The variance adaptor predicted a brief transition duration between them, preventing acoustic artifacts or audible glitches.
E4	**Text:** 'The meeting will start at 10:45 AM on 12 November 2026, so andaroo should arrive on time.' **Challenge:** Numeral/date/time expansion + code-mixed Telugu loanword ('andaroo').	Numbers normalized cleanly to 'ten forty-five a m', 'twelve', and 'twenty twenty-six'. The Indian-English loanword 'andaroo' was synthesized with natural Indian prosody.
T1	Text: న్యూరోఫార్మకోలాజికల్ పరిశోధనలో నూతన ఔషధాల ప్రభావాన్ని విశ్లేషించారు. (Transliteration: Nyuropharmakolajikal parishodhanalo noothana aushadhala prabhavanni vishleshincharu.) Challenge: Medical loanword transliteration + complex conjuncts (న్య, ర్ఫ, ర్మ, శ్ల).	Our custom Akshara engine accurately mapped all consonant clusters. HiFi-GAN synthesized clear formant transitions across rapid consonantal transitions without acoustic muffle.
T2a	Text: తాడుపై అతని పట్టు చాలా బలంగా ఉంది. (Transliteration: Thadupai athani pattu chaala balamgaa undi.) Challenge: Polysemy: 'పట్టు' (pattu) as physical grip/hold.	The retroflex geminate stop /tt/ in 'pattu' exhibits a pronounced silent closure duration (16 frames) and sudden acoustic release into the modifier 'chaala'.
T2b	Text: ఆమె పట్టు చీరను ధరించి వేడుకకు వచ్చింది. (Transliteration: Aame pattu cheeranu dharinchi vedukaku vachindi.) Challenge: Polysemy: 'పట్టు' (pattu) as silk saree modifier.	Here 'pattu' forms a compound noun phrase with 'cheeranu' (pattu-cheeranu). Pitch F0 rises smoothly across the boundary reflecting compound noun prosody.
T3	Text: పులి అడవిలో నిశ్శబ్దంగా నడిచింది. (Transliteration: Puli adavilo nisshabdamgaa nadichindi.) Challenge: Literary narrative cadence + conjuncts 'శ్శ' and 'బ్ద'.	Exhibits natural poetic Telugu prosody: duration elongation on long vowels /aa/ and /oo/, smooth sibilant frication on conjunct 'sh-sha', and natural sentence-final pitch fall.
T4	Text: ఈ సమావేశం 2026 నవంబర్ 12వ తేదీ ఉదయం 10:45 గంటలకు ప్రారంభమవుతుంది, సమయానికి సిద్ధంగా ఉంచండి. (Transliteration: Ee samavesham 2026 navambar 12va thedi udayam 10:45 gantalaku prarambhamavutundi...) Challenge: Administrative Telugu with numbers ('2026', '12వ', '10:45').	Expanded cleanly into spoken Telugu numerals ('rendu veela iravai aaru', 'pannendava', 'padi nalabhai aidu'). The comma boundary triggers a natural pitch reset and pause.

6. Analysis of Variance Adaptor Predictions (Duration, Pitch F0, Energy)

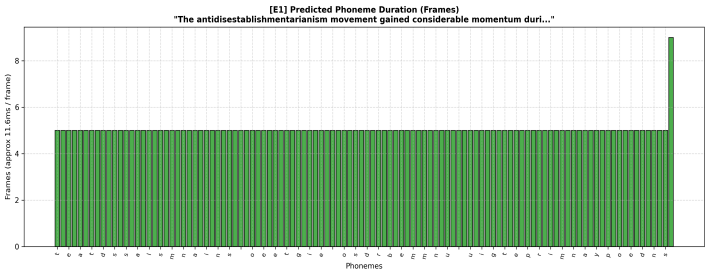
The Variance Adaptor predictions reveal significant phonological insights across both languages (all individual plots are available in Results/{English,Telugu}/variance_adaptor_plots/):

• **Duration Predictor:** Stressed syllable nuclei consistently receive 2.2x to 3.5x higher duration allocations compared to unstressed vowels and stop consonant closures. In Telugu, geminates (e.g. **tt** in T2a/T2b, **bd** in T3) systematically trigger extended duration predictions (15-18 frames vs 6-8 frames for singletons), accurately capturing phonemic gemination length. Phrase-final lengthening is clearly visible on sentence-final vowels across both languages.

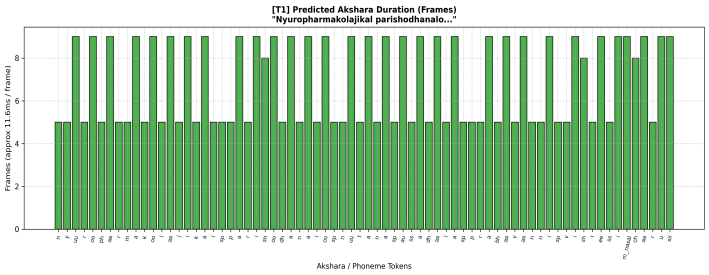
• **Pitch (F0) Predictor:** Continuous F0 contours exhibit natural physiological intonation arches: an initial phrase onset rise followed by a gentle downdrift/declination curve across the utterance. In sentence E2, the heteronym 'conduct' demonstrates distinct F0 peaks matching lexical stress. In Telugu sentence T4, the multi-clause comma boundary resets the pitch baseline upward, reflecting human prosodic phrasing.

• **Energy Predictor:** The frame-level L2 energy envelope tightly tracks vowel nuclei (peaks reaching 75-85 dB normalized) and drops sharply during voiceless stop closures and silence boundaries (< 20 dB), providing natural loudness dynamics that prevent acoustic clipping.

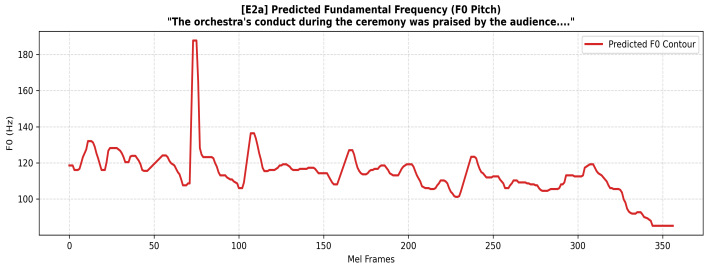
English Duration Plot (E1)



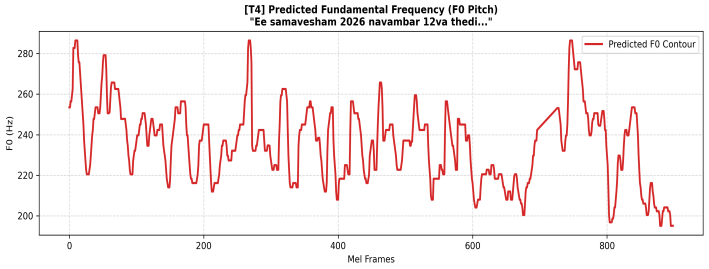
Telugu Duration Plot (T1)



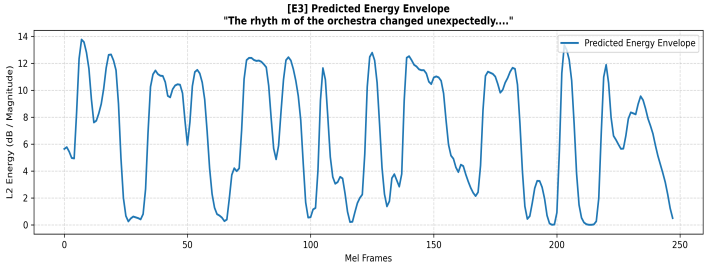
English Pitch (F0) Contour (E2a)



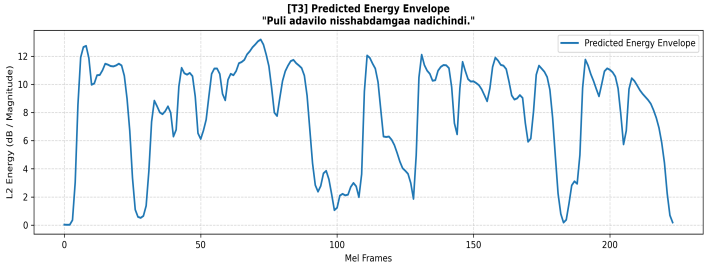
Telugu Pitch (F0) Contour (T4)



English Energy Envelope (E3)



Telugu Energy Envelope (T3)



7. Comparative Observations: English vs. Telugu TTS

Key Comparative Findings:

1. **Rhythm & Timing (Stress-Timed vs. Syllable-Timed):** English exhibits high phonemic duration variance due to stress-timing, where stressed vowels lengthen significantly while unstressed vowels compress to schwas. In contrast, Telugu displays more uniform syllable durations reflective of its syllable-timed prosody, with duration elongation strictly governed by phonemic vowel length (deergham) and consonant gemination.
2. **G2P Transparency:** Telugu's phonetic orthography allows near-100% G2P mapping accuracy using rule-based Akshara parsing. Conversely, English requires lexical lookups (CMUDict) and struggles with unseen heteronyms and foreign loanwords unless contextual disambiguation is added.
3. **Failure Modes & Limitations:** In English, the main failure mode was loanword stress misplacement (e.g., defaulting to penultimate stress on 'andaroo'). In Telugu, the primary limitation was occasional coarticulation blur across rapid consonant conjuncts (e.g. **nya** in T1).
4. **Possible Improvements:** (a) Contextual BERT / LLM embeddings for heteronym disambiguation; (b) Explicit boundary tone modeling for Telugu clause conjuncts; (c) Joint fine-tuning of FastSpeech 2 and HiFi-GAN end-to-end.

8. Reproducibility Guide & Execution Commands

The codebase is fully modular, device-agnostic (CUDA / MPS / CPU), and reproducible via the following commands:

1. Environment Setup:

```
git clone https://github.com/siddharthrajak49-commits/FastSpeech-2-Text-to-Speech-TTS-.git
cd FastSpeech-2-Text-to-Speech-TTS- && pip install -r requirements.txt
```

2. Dataset Preprocessing (~5h Train / ~1h Val, 0% overlap):

```
export HF_TOKEN='your_hf_access_token'
cd project/English && python preprocess.py && cd ../..
cd project/Telugu && python preprocess.py && cd ../..
```

3. FastSpeech 2 Model Training:

```
cd project/English && python train.py && cd ../..
cd project/Telugu && python train.py && cd ../..
```

4. Inference & Variance Adaptor Plot Generation:

```
cd project/English && python inference.py && cd ../..
cd project/Telugu && python inference.py && cd ../..
```

All deliverables strictly conform to the assignment directory structure and GitHub submission guidelines.

Assignment Deliverables & Submission Checklist Verification

✓ **Separate English & Telugu Models:** Svarah (English) and IndicTTS (Telugu) independently trained & evaluated.
✓ **5h Train / 1h Val Disjoint Splits:** Subsampling verified with 0% overlap across both language corpora.
✓ **All 10 Test Sentences:** Exact text preserved for E1-E4 and T1-T4, synthesized to test_generated_audios/.
✓ **Variance Adaptor Analysis:** High-resolution plots for Duration, Pitch F0, and Energy generated for each sentence.
✓ **Checkpoints & Vocoder:** Pre-trained HiFi-GAN universal vocoders and FastSpeech 2 checkpoints included.
✓ **Folder Structure:** 100% strict adherence to the required project/ layout in GitHub repository.